

# Configuration of a linear solver for linearly implicit time integration and efficient data transfer in parallel thermo-hydraulic computations

Ravil Dorozhinskii

Computational Science and Engineering

March 21, 2019

# Outline I

## 1 Overview of ATHLET and NuT software

- ATHLET
- NuT

## 2 Problem Statement

## 3 Methodology and Matrix Sets

## 4 Overview of Solver Types

- Iterative Methods
  - Theory overview
  - Aspects
- Direct Sparse Methods
  - Theory overview
  - Parallelization Aspects
  - Numerical Accuracy
- Conclusion

# Overview of ATHLET and NuT software

# Mathematical Model

## 1. Liquid mass

$$\frac{\partial((1 - \alpha)\rho_l)}{\partial t} + \nabla((1 - \alpha)\rho_l \vec{w}_l) = -\psi \quad (1)$$

## 2. Vapor mass

$$\frac{\partial(\alpha\rho_v)}{\partial t} + \nabla(\alpha\rho_v \vec{w}_v) = \psi \quad (2)$$

## 3. Liquid momentum

$$\frac{\partial((1 - \alpha)\rho_l \vec{w}_l)}{\partial t} + \nabla((1 - \alpha)\rho_l \vec{w}_l \vec{w}_l) + \nabla((1 - \alpha)p) = \vec{F}_l \quad (3)$$

## 4. Vapor momentum

$$\frac{\partial(\alpha\rho_v \vec{w}_v)}{\partial t} + \nabla(\alpha\rho_v \vec{w}_v \vec{w}_v) + \nabla(\alpha p) = \vec{F}_v \quad (4)$$

# Mathematical Model

## 5. Liquid energy

$$\frac{\partial \left[ (1 - \alpha) \rho_l \left( h_l + \frac{1}{2} \vec{w}_l \vec{w}_l - \frac{p}{\rho_l} \right) \right]}{\partial t} + \nabla \left[ (1 - \alpha) \rho_l \vec{w}_l \left( h_l + \frac{1}{2} \vec{w}_l \vec{w}_l \right) \right] = -p \frac{\partial (1 - \alpha)}{\partial t} + E_l \quad (5)$$

## 6. Vapor energy


$$\frac{\partial \left[ \alpha \rho_v \left( h_v + \frac{1}{2} \vec{w}_v \vec{w}_v - \frac{p}{\rho_v} \right) \right]}{\partial t} + \nabla \left[ \alpha \rho_v \vec{w}_v \left( h_v + \frac{1}{2} \vec{w}_v \vec{w}_v \right) \right] = -p \frac{\partial \alpha}{\partial t} + E_v \quad (6)$$

## 7. Volume vapor fraction

$$\alpha = \frac{V_v}{V} \quad (7)$$

According to [3]

# Mathematical Model: from PDEs to ODEs



**Figure 1:** ATHLET: one dimensional **Finite Volume** formulation of the problem [13]

The system is transformed to a non-autonomous system of **ordinary differential equations** and expressed as an **initial value problem** after spatial finite-volume integration and some mathematical transformations [3].

$$\frac{dy}{dt} = f(t, y), \quad t_0 \leq t \leq t_F \quad y(t_0) = y_0 \quad (8)$$

where  $y \in \mathbb{R}^N$  is a composite vector of variables,  $f$  is a non-linear function such that  $f : \mathbb{R} \times \mathbb{R}^N \supset \Omega \rightarrow \mathbb{R}^N$ .

# Numerical Integration: 6-stage W-method

$$(I - h_i J) \delta y_{ij} = h_i f(t_0 + j \cdot h_i, y_0 + \sum_{l=0}^{j-1} \delta y_{il}) + h^2 \frac{\partial f_0}{\partial t} \quad (9)$$

where  $j = 0, \dots, i-1$ ;  $h_i = h/i$ ;  $i = 1, 2, 3$

Implicit Euler with one Newton's iteration




Figure 2: A general view on 6-stage W-method [13]

# Problem Statement



# Solving Sparse Linear Systems

Numerical **integration** of a system of ODEs by means of **W-methods**:

$$(I - h_i J) \delta y_{ij} = h_i f(t_0 + j \cdot h_i, y_0 + \sum_{l=0}^{j-1} \delta y_{il}) + h^2 \frac{\partial f_0}{\partial t} \quad (10)$$

can be considered as a solution of a **sequence of linear systems** from another point of view:

$$A_i \delta y_{ij} = b_{ij} \quad (11)$$

where  $A = (I - h_i J)$  is a  $\mathbb{R}^N \times \mathbb{R}^N$  non-singular **sparse matrix**;  $\delta y_{ij}$  and  $b_{ij}$  are  $\mathbb{R}^N$  vectors;  $i = 1, 2, 3; j = 0, \dots, i - 1$

## NOTE

Computational **burden** of the W-method mainly lies in **solving of sparse linear systems** beside evaluation of non-linear function  $f$  of Eq. 8

# Methodology and Matrix Sets

# Static Solver Configuration

- ATHLET is a CFD tool dedicated to **transient problems**
- Additionally, **topology** of hydrolic circuits can **be changed** during simulation-time
- Hence, the Jacobian matrix **structure** can **vary** significantly during numerical itegration

## Therefore:

- 1 Configuration of a linear solver in run-time is compute-intensive and time consuming
- 2 Reults of dynamic solver configuration can be ambiguous and difficult to interpret
- 3 A feasible and doable approach is **static solver configuration**

# GRS Matrix Set

GRS **matrix set** was generated by **running** the most **common GRS simulations** in ATHLET and **stopping** them **somewhere in the middle**. The corresponding shifted Jacobian matrices were saved in the PETSc binary format.

Name	n	nnz	nnz / n	Approximate Condition Number
pwr-3d	6009	32537	5.4147	1.019e+07
cube-5	9325	117897	12.6431	1.592e+09
cube-64	100657	1388993	13.7993	7.406e+08
cube-645	1000045	13906057	13.9054	6.474e+08
k3-2	130101	787997	6.0568	1.965e+15
k3-18	1155955	7204723	6.2327	1.947e+12

Table 1: GRS matrix set

# SuiteSparse Matrix Set

**SiteSparse** matrix set was generated by **downloading** a dozen of matrices from SuiteSparse Matrix Collection [5], [6], with the **aim** of **comparison** and **verification** of solver configurations

Name	n	nnz	nnz / n	Approximate Condition Number
cant	62451	4007383	64.1684	5.082e+05
consph	83334	6010480	72.1251	2.438e+05
CurlCurl.3	1219574	13544618	11.1060	2.105e+05
Geo_1438	1437960	63156690	43.9210	4.677e+05
memchip	2707524	13343948	4.9285	1.305e+07
PFlow_742	742793	37138461	49.9984	5.553e+06
pkustk10	80676	4308984	53.4110	5.589e+02
torso3	259156	4429042	7.0903	2.456e+03
x104	108384	8713602	80.3956	3.124e+05

Table 2: SuiteSparse matrix set

# Hardware

	HW1 (GRS)	HW2 (LRZ Linux)
CPU(s)	20	28
On-line CPU(s) list	0-19	0-27
Thread(s) per core	1	1
Core(s) per socket	10	14
Socket(s)	2	2
NUMA node(s)	2	4
Model name	E5-2680 v2	E5-2697 v3
Stepping	4	2
CPU MHz	1200.0	2036.707
L1 d/i cache	32K/32K	32K/32K
L2 cache	256K	256K
L3 cache	25600K	17920K
NUMA node0 CPU(s)	0-9	0-6
NUMA node1 CPU(s)	10-19	7-13
NUMA node2 CPU(s)	-	14-20
NUMA node3 CPU(s)	-	21-27

Table 3: Hardware specification

# Experimental setup

## Libraries and Compiler

- 1 PETSc version 3.10
- 2 OpenMPI version 3.1.1
- 3 Intel Compiler 18

# Overview of Solver Types



# Iterative methods

- Given an initial guess
- An iterative method generates a sequence of approximate solutions by means of a specific rule
- Depending on a method and the given problem,
- there may exist certain conditions such that the sequence eventually converges to the exact solution
- There exist two families of iterative methods: *stationary* and *Krylov-based* methods
- Nowadays, Krylov methods dominate in the field of scientific computing
  - because of their rather fast convergence
  - in case of solving well conditioned systems
  - or/and a "good" initial guess

# Krylov-based methods

- *The key idea* is to construct an approximate solution
- as a linear combination of vectors  $b, Ab, A^2b, A^3b, \dots A^{n-1}b$
- known as Krylov subspace  $\mathcal{K}_n$
- where, without loss of generality, the initial guess  $x_0$  is equal to zero
- At each iteration, the subspace is expanded by adding and evaluating the next vector in the sequence
- the methods define and expand another subspace  $\mathcal{L}_n$
- such that  $r_n = b - Ax_n \perp \mathcal{L}_n$
- which is known as the Petrov-Galerkin condition
- A construction of subspace  $\mathcal{L}_n$  is defined by a method
- and based on matrix properties

# Aspects

## Parallelization Aspects

- Iterative methods usually make use of simple linear algebra kernels e.g. dot products, matrix-vector products, etc.
- Therefore, the methods efficiently make use of data-based parallelism

## Numerical Accuracy

- Convergence to the exact solution depends on a value of the condition number of a matrix
- Usually requires a linear transformation known as preconditioning
  - to reduce condition number and accelerate convergence

# Preconditioners

Package name	Origin	Method	Tuning parameters	Comments
block Jacobi	PETSc	block Jacobi	-pc.bjacobi.blocks -sub_pc_type	-
additive Schwarz	PETSc	additive Schwarz	-pc.asm.blocks -pc.asm.overlap -pc.asm.type -pc.asm.local_type -sub_pc_type	-
euclid	hypre	ILU(k)	-nlevel -thresh -filter	deprecated form PETSc
pilut	hypre	ILU(t)	-pc.hypre.pilut_tol -pc.hypre.pilut_maxiter -pc.hypre.pilut_factorrowsize	-
parasail	hypre	SPAI	-pc.hypre.parasails_nlevels -pc.hypre.parasails_thresh -pc.hypre.parasails_filter	-
SPAI	Grote, Barnard	SPAI	-pc.spai_epsilon -pc.spai_nbstep -pc.spai_max -pc.spai_max_new -pc.spai_block_size -pc.spai_cache_size	-
BoomerAMG	hypre	algebraic multigrid	-pc.hypre.boomeramg_cycle_type -pc.hypre.boomeramg_max_levels -pc.hypre.boomeramg_max_iter -pc.hypre.boomeramg_tol etc.	39 tuning parameters in total

Table 4: Parallel preconditioning algorithms available in PETSC

# Direct Sparse Methods

Direct sparse methods combine the main advantages of direct and iterative methods

- numerical accuracy of the methods is comparable with the standard Gaussian Elimination process
- complexity is bounded by  $O(n^2)$  due to efficient treatment of sparsity

A solution of a system of equations is computed:

- by means of forward and backward substitutions
- using  $LU$  decomposition of the corresponding matrix.

# Direct Sparse Methods: Speech

- The multifrontal method is probably the most representative example of direct sparse solvers introduced by Duff and Reid
- The method is an improved version of the frontal method
- which can compute independent fronts in parallel.
- A front, also called a frontal matrix, can be considered as a small dense matrix
- resulting from a column elimination of the original system.
- There also exist left- and right-looking variants of the multifrontal method

Let's consider the multifrontal methods as an example

To keep the overview rather simple, we assume that matrix  $A$  is **symmetric positive definite** and **sparse**

# Multifrontal Method: I

$$A = LDL^T \quad \text{with} \quad (D)_{ii} > 0 \quad (12)$$

Given a rule and a sparsity pattern, one can build an elimination tree

$$p = \min(i > j | l_{ij} \neq 0) \quad (13)$$

$$A = \begin{pmatrix} 1 & a & & & & & & & \\ 2 & & b & & \bullet & & & & \\ 3 & & & c & & \bullet & & & \\ 4 & & \bullet & & d & & & & \\ 5 & & & \bullet & & e & & & \\ 6 & & \bullet & & & & f & & \\ 7 & \bullet & & & & & & g & \\ 8 & \bullet & \bullet & \bullet & & & & & h \\ 9 & \bullet & & \bullet & \bullet & & & & & i \end{pmatrix} \quad L = \begin{pmatrix} 1 & a & & & & & & & \\ 2 & & b & & & & & & \\ 3 & & & c & & & & & \\ 4 & & \bullet & & d & & & & \\ 5 & & & \bullet & & e & & & \\ 6 & & \bullet & & & & f & & \\ 7 & \bullet & & & & & & g & \\ 8 & \bullet & \bullet & \bullet & & & & & h \\ 9 & \bullet & & \bullet & \bullet & & & & & i \end{pmatrix}$$




Figure 3: A sparse matrix and its Cholesky factor, [10]

Figure 4: An elimination tree, [10]

## Multifrontal Method: II

The fundamental idea of the multifrontal method spins around frontal  $F$  and update matrices  $\hat{U}$

$$F_j = Fr_j + \hat{U}_j = \begin{bmatrix} a_{j,j} & a_{j,i_1} & a_{j,i_2} & \dots & a_{j,i_r} \\ a_{i_1,j} & & & & \\ a_{i_2,j} & & & & \\ \vdots & & & 0 & \\ a_{i_r,j} & & & & \end{bmatrix} + \hat{U}_j \quad (14)$$

where  $i_0, i_1, i_2, \dots, i_r$  are row subscripts of non-zeros in  $L_{*j}$  where  $i_0 = j$ ;  $r$  is the number of off-diagonal non-zero elements.

$$\hat{U}_j = - \sum_{k \in T[j]-j} \begin{bmatrix} l_{j,k} \\ l_{i_1,k} \\ \vdots \\ l_{i_r,k} \end{bmatrix} [l_{j,k} \quad l_{i_1,k} \quad \dots \quad l_{i_r,k}] \quad (15)$$

where  $\hat{U}_j$  can be treated as the second term of the Schur complement



## Multifrontal Method: III

Let's consider factorization of a 2-by-2 block matrix  $A$

$$A = \begin{bmatrix} B & V^T \\ V & C \end{bmatrix} = \begin{bmatrix} L_B & 0 \\ VL_B^{-T} & I \end{bmatrix} \begin{bmatrix} I & 0 \\ 0 & Sr \end{bmatrix} \begin{bmatrix} L_B^T & L_B^{-1}V^T \\ 0 & I \end{bmatrix} \quad (16)$$

Assuming that block  $B$  has already been factorized:  $B = L_B L_B^T$

The Schur complement can be viewed as:

$$Sr = C - VB^{-1}V^T \quad (17)$$

*Dense Linear Algebra:*  $-VB^{-1}V^T$

*Multifrontal method:*  $\hat{U}_j$

$$-\sum_{k=1}^{j-1} \begin{bmatrix} l_{j,k} \\ \vdots \\ l_{n,k} \end{bmatrix} [l_{j,k} \quad \dots \quad l_{n,k}] \quad (18)$$

$$-\sum_{k \in T[j]-j} \begin{bmatrix} l_{i_0,k} \\ \vdots \\ l_{i_r,k} \end{bmatrix} [l_{i_0,k} \quad \dots \quad l_{i_r,k}] \quad (19)$$

# Multifrontal Method: IV

After column  $j$  factorization:

$$\hat{F}_j = \begin{bmatrix} l_{j,j} & \dots & 0 \\ \vdots & I & \\ l_{i_r,j} & & \end{bmatrix} \begin{bmatrix} 1 & \dots & 0 \\ \vdots & U_j & \\ 0 & & \end{bmatrix} \begin{bmatrix} l_{j,j} & \dots & l_{i_r,j} \\ \vdots & I & \\ 0 & & \end{bmatrix} \quad (20)$$

where sub-matrix  $U_j$  represents the full update from all descendants of node  $j$  and node  $j$  itself




Figure 5:  
Information flow

(21)

# Supernodal Method

- In practice, an improved version of the multifrontal method is used
- A super-node is formed by a set of contiguous columns
- which have the same off-diagonal sparsity structure

$$\mathcal{F}_j = \begin{bmatrix} a_{j,j} & a_{j,j+1} & \dots & a_{j,j+t} & a_{j,i_1} & \dots & a_{j,i_r} \\ a_{j+1,j} & a_{j+1,j+1} & \dots & a_{j+1,j+t} & a_{j+1,i_1} & \dots & a_{j+1,i_r} \\ \vdots & \vdots & \dots & \vdots & & & \\ a_{j+t,j} & a_{j+t,j+1} & \dots & a_{j+t,j+t} & a_{j+t,i_1} & \dots & a_{j+t,i_r} \\ a_{i_1,j} & a_{i_1,j+1} & \dots & a_{i_1,j+t} & & & \\ \vdots & \vdots & \dots & \vdots & & 0 & \\ a_{i_r,j} & a_{i_r,j+1} & \dots & a_{i_r,j+t} & & & \end{bmatrix} \quad \begin{matrix} \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \end{matrix} U_{C_1} \dots \begin{matrix} \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \\ \text{\texttt{+}} \end{matrix} U_{C_s}$$

(22)

# Parallelization Aspects




Figure 6: Original elimination tree




Figure 7: Supernodal/assembly tree

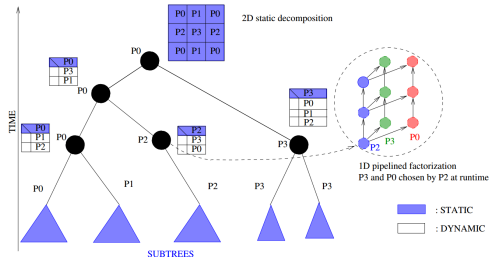


Figure 8: Static and dynamic scheduling in MUMPS, [8]

# Numerical Accuracy: I

- Pivoting is a big issue in direct sparse methods
  - analysis: absence of numerical information
  - numerical factorization: it may distort all prediction made during the analysis phases i.e.
    - fill-in prediction
    - load balancing

Therefore, threshold pivoting is commonly used for direct sparse methods

Threshold pivoting means that a pivot  $|a_{i,i}|$  is accepted if it satisfies:

$$|a_{i,i}| \geq \alpha \times \max_{k=i \dots n} |a_{k,i}| \quad (23)$$

where  $\alpha \in [0, 1]$  and  $k = i \dots n$  represents row indices of column  $i$  within the fully summed block of a frontal matrix.

## Numerical Accuracy: II

- In case of small values of  $\alpha$
- solutions can be numerically inaccurate
- may demand to perform solution refinements

As an example, solution accuracy can be improved using:

- iterative refinement method based a on solution residual
- resulting  $LU$  decomposition can be used as a preconditioner for a Krylov-based method e.g. GMRES

# Results and Conclusion

- As the first step, various preconditioning algorithms were tested
  - A coarse grid search was used
  - with maximum 3 values for each tuning parameter
  - starting from the default towards more accurate values
  - Testing showed there was no preconditioning algorithm that could result in convergence for the entire set of matrices
- Even if we can find an algorithm and suitable parameter settings
  - which will result in convergence of the entire matrix set
  - there is no guarantee that it will work in all steps during a simulation
  - Iterative methods cannot fulfill *robustness* criterion
- Therefore, direct sparse methods is only one way to go
- In spite of limited tree-task parallelism

# Overview of available Direct Sparse Solver (DSS) libraries

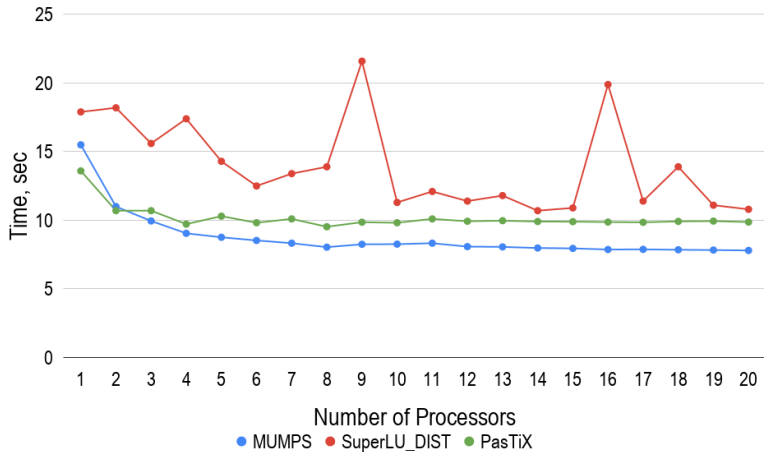


# List of parallel DSSs

Package	Method	Matrix Types	PETSc Interface	License
Clique	Multifrontal	Symmetric	Not Officially	Open
MF2	Multifrontal	Symmetric pattern	No	-
DSCPACK	Multifrontal	SPD	No	Open
MUMPS	Multifrontal	General	Yes	Open
PaStiX	Left looking	General	Yes	Open
PSPASES	Multifrontal	SPD	No	Open
SPOOLES	Left-looking	Symmetric pattern	No	Open
SuperLU_DIST	Right-looking	General	Yes	Open
symPACK	Left-Right looking	SPD	No	Open
S+	Right-lookin	General	No	-
PARDISO	Multifrontal	General	No	Commercial
WSMP	Multifrontal	General	No	Commercial

**Table 5:** A list of direct sparse linear solvers adapted for distributed-memory computations, [9], [4]

# Comparisons of DSSs



**Figure 9:** Comparisons of MUMPS, PasTiX and SuperLU\_DIST libraries during 5 point-stencil Poisson matrix (1000000 equations) factorizations

## Configuration of MUMPS solver

# ABCD

# Fill Reducing Reorderings

## Sequential:


- Approximate Minimum Degree (AMD) [2]
- Approximate Minimum Fill (AMF)
- Approximate Minimum Degree with automatic quasi-dense row detection (QAMD) [1]
- Bottom-up and Top-down Sparse Reordering (PORD) [12]
- Nested Dissection coupled with AMD (Scotch) [11]
- Multilevel Nested Dissection coupled with Multiple Minimum Degree (METIS) [7]

## Parallel:


- ParMETIS
- PT-Scotch

# Explicit MPI Process Pinning

To make process pinning **deterministic**, a python script was developed to automatically generate **rankfiles** based on the number of **MPI processes**, **OpenMP threads** per MPI process, the maximum number of processing elements and the number of **NUMA domains**.



(a) Spread



(b) Close

Figure 10: Process pinning of 2 OpenMP / 1 MPI on HW1 hardware

# ABCD

# ABCD



# ABCD

# ABCD

# Outline II

## 5 Overview of available Direct Sparse Solver (DSS) libraries

## 6 Configuration of MUMPS solver

- Overview of MUMPS
- Fill Reducing Reorderings
- Process pinning
- Configuration of BLAS library
- Hybrid MPI-OpenMP process/thread distribution
- Conclusion
- Recommendations

*Thanks for your attention!*



PR Amestoy. “Recent progress in parallel multifrontal solvers for unsymmetric sparse matrices”. In: *Proceedings of the 15th World Congress on Scientific Computation, Modelling and Applied Mathematics, IMACS*. Vol. 97. 1997.



Patrick R. Amestoy, Timothy A. Davis, and Iain S. Duff. “An Approximate Minimum Degree Ordering Algorithm”. In: *SIAM Journal on Matrix Analysis and Applications* (1996).



H. Austregesilo et al. *ATHLET Mod 3.1A – Models and Methods*. distributed with ATHLET. Mar. 2016.



Satish Balay et al. *PETSc Web page*.  
<http://www.mcs.anl.gov/petsc>. 2018. URL:  
<http://www.mcs.anl.gov/petsc>.



Timothy A. Davis and Yifan Hu. “The University of Florida Sparse Matrix Collection”. In: *ACM Trans. Math. Softw.* 38.1 (Dec. 2011), 1:1–1:25. ISSN: 0098-3500. DOI: 10.1145/2049662.2049663. URL: <http://doi.acm.org/10.1145/2049662.2049663>.



I. S. Duff, Roger G. Grimes, and John G. Lewis. “Sparse Matrix Test Problems”. In: *ACM Trans. Math. Softw.* 15.1 (Mar. 1989), pp. 1–14. ISSN: 0098-3500. DOI: 10.1145/62038.62043. URL: <http://doi.acm.org/10.1145/62038.62043>.



George Karypis and Vipin Kumar. *MeTis: Unstructured Graph Partitioning and Sparse Matrix Ordering System, Version 4.0*. <http://www.cs.umn.edu/~metis>. University of Minnesota, Minneapolis, MN, 2009.



Jean-Yves L'Excellent. “Multifrontal methods: parallelism, memory usage and numerical aspects”. PhD thesis. Ecole normale supérieure de lyon-ENS LYON, 2012.



X. Li. *Direct Solvers for Sparse Matrices*. 2018. URL: <http://crd-legacy.lbl.gov/~xiaoye/SuperLU/SparseDirectSurvey.pdf>.



Joseph W. H. Liu. “The Multifrontal Method for Sparse Matrix Solution: Theory and Practice”. In: *SIAM* (1992).



François Pellegrini. “Scotch and libScotch 5.1 user’s guide”. In: (2008).



Jürgen Schulze. “Towards a tighter coupling of bottom-up and top-down sparse matrix ordering methods”. In: *BIT Numerical Mathematics* 41.4 (2001), pp. 800–841.



Tim Steinhoff. “Singly implicit FiterRK methods for thermal-hydraulic simulations”. 2018. URL: <http://somewhere.in.grs.com>.